

AddLyft · Social Post Workflow

A plain-language overview of how text posts and poster images are created
Client overview · 23 Sep 2026 · Social post only (video not included)

How to read numbers in this document

Measured = taken from our project usage logs · **Configured estimate** = calculated with the price rates coded in the app (helpful for planning; not a Google Cloud invoice) · **Unknown** = not recorded yet in this demo.

1. In one minute

Someone on your team types a short promotion idea (the “brief”). The store is already set for this demo as **AMPM Woodstock**, so they do not need to repeat the store name every time. They can optionally upload up to three product or style photos.

One click creates **two drafts to review**: social-media wording (caption, offer, call-to-action, hashtags) and a vertical poster image. Nothing is posted automatically — a person still decides what goes live.

Topic	Detail
Primary screen	Local showcase at http://127.0.0.1:5173/
Text AI	Gemini 2.5 Flash-Lite (Google Vertex AI)
Image AI	Nano Banana lite image model (9:16 portrait, 1K)
Typical full generate cost	About \$0.0345 (Configured estimate from Measured logs: ~\$0.0001 text + ~\$0.0344 image)
Typical poster wait time	About 8–11.5 seconds (Measured from recent runs)
Text wait time	Unknown in logs (not recorded for text yet)

2. The big picture

Text writing and image creation run **at the same time** (in parallel). If one fails, the other can still succeed — you keep whatever worked.




Figure A — One brief feeds two drafts; you review before anything is published.

How many AI calls? A normal Generate uses **2** model calls (one text, one image). “Regenerate” updates wording only (1 text call). The image path does not auto-edit; a correction would be a separate step.

3. Text workflow (social copy)

The text path turns your brief into structured post fields that look like a finished social post in the preview — not raw JSON for the client to decode.



Figure B — Text path, step by step.

What you get: caption, offer/key message, CTA button text, and 3–6 hashtags. The selected store is added behind the scenes so the copy can mention AMPM Woodstock without you typing it in the brief.

4. Image workflow (poster)

The image path builds a detailed creative brief for the image model, optionally using your photos as product or style guides, then returns one portrait poster.




Figure C — Image path, step by step.

Topic	Detail
Format	Portrait 9:16 · 1K (lite model limit)
Photos	0–3 optional · roles: product / style / brand / background / auto
Photo prep	Resized to max side 1600px JPEG before send
People in the poster	Off by default · turn on only if you want faces (Measured policy in product)
API calls per generate	Exactly 1 image model call (happy path)

5. Cost, tokens, and timing

Google bills AI usage mainly by how much text/image “work” each run uses (tokens). Longer briefs and reference photos usually mean more tokens. Our app shows estimated cost after each run using rates we configured for planning — useful for R&D budgeting, not a substitute for the official cloud bill.

5.1 Price rates used in the estimator

Part	What it covers	USD / 1M tokens	Kind
Text input	Brief + instructions	\$0.10	Configured estimate
Text output	Generated post fields	\$0.40	Configured estimate
Image input	Prompt + photo tokens	\$0.25	Configured estimate
Image pixels	Generated poster	\$30.00	Configured estimate

For a 1K poster, the estimator often counts about **1120** image-output tokens when the provider does not return a finer split (Configured table in code).

5.2 Recent text runs (from logs)

When (UTC)	Tokens in / out / total	Est. USD	Wait time
2026-09-22 19:41	367 / 156 / 523	\$0.000099	Unknown
2026-09-22 19:33	367 / 133 / 500	\$0.000090	Unknown

2026-09-13 13:54	47 / 82 / 129	\$0.000038	Unknown
------------------	---------------	------------	---------

Measured token counts · Configured estimate dollars · model: gemini-2.5-flash-lite. Text wait time is Unknown (not stored in the text log).

5.3 Recent poster runs (from logs)

When (UTC)	Total tokens	Est. USD	Wait (seconds)
2026-09-22 19:41	4417	\$0.03442	11.43
2026-09-21 18:40	4360	\$0.03441	11.36
2026-09-21 18:37	4360	\$0.03441	11.07
2026-09-21 18:31	1716	\$0.03375	8.04

Measured · model: gemini-3.1-flash-lite-image · 9:16 · 1K. Poster cost dominates a full Generate (over 99% of the combined estimate).

5.4 Takeaway for budgeting

Plan on roughly **three to four cents per full text+poster generate** at current lite settings (Configured estimate from Measured logs). Wording-only regenerates are far cheaper (fractions of a cent). Cloud discounts or rate changes can move the real invoice.

6. Guardrails — what the system protects

We do not rely on the AI “being careful” alone. Checks happen before and after generation, and the UI makes it clear these are drafts.



The app proposes drafts. A person still decides what goes live.

Figure D — Four layers from input to human approval.

- **Input:** empty briefs rejected; 2000-character limit; simple blocklist for clearly abusive asks; max 3 photos.
- **During AI:** Google safety filters for hate, harassment, sexual, and dangerous content; “facts only” instructions (no invented discounts or dates); store injected from selection.
- **After AI (text):** required fields, length limits, hashtag formatting, light check that \$ / % style claims appeared in the brief.
- **You:** on-screen “draft — review before posting” checklist. Legal proof for strong claims (e.g. “#1”) remains a business decision.

7. What this demo does not do (yet)

- It does not post to Instagram, Facebook, TikTok, or other networks.
- It is not a full multi-store app — store is fixed for the demo (placeholder for a future “pick your store” step).
- It does not automatically proofread every letter on the poster (no OCR quality gate yet).
- Estimated dollars are not the final Google Cloud invoice.
- Ad video / Veo is out of scope for this social-post package.

8. Technical snapshot (for your team)

Useful if engineering wants exact names. Clients can skip this section.

Topic	Detail
UI	React (Vite) · port 5173 · proxies /api → 8000
API	FastAPI · POST /api/generate · POST /api/generate-image
Text model ID	gemini-2.5-flash-lite · Vertex · us-central1 (default)
Image model ID	gemini-3.1-flash-lite-image · Vertex · global · 9:16 · 1K

Logs used for Measured rows	logs/usage.jsonl · social-post-image/logs/image_usage.jsonl
Related PDFs	social_post_generation_report.pdf · AddLyft_Social_Post_Image_MVP_Audit.pdf (earlier image-only view)